

Assignment 3.1: Group Comparison

Import and Install

```
In [1]: import os
import time
import re
import random
import json
from pathlib import Path
from urllib.parse import urljoin
from collections import Counter

import requests
import pandas as pd
from bs4 import BeautifulSoup
from slugify import slugify

import matplotlib.pyplot as plt
from wordcloud import WordCloud

import nltk
nltk.download("punkt", quiet=True)
from nltk.tokenize import word_tokenize
```

Configure

```
In [2]: BASE = Path(r"C:\Users\gabed\ADS-509-Assignment-2.1\data\M1 Results")
TWITTER_DIR = BASE / "twitter"
LYRICS_DIR = BASE / "lyrics" # will be created if missing
FIG_DIR = Path("./figures"); FIG_DIR.mkdir(exist_ok=True)

# Your specific files (from your screenshots) - include the .txt extension
TWITTER_FILES = {
    "the_who": TWITTER_DIR / "The_Who_followers_data.txt",
    "rush": TWITTER_DIR / "Rush_followers_data.txt",
}

ARTISTS_URLS = {
    "the_who": "https://www.azlyrics.com/w/who.html",
    "rush": "https://www.azlyrics.com/r/rush.html",
}

MAX_SONGS_PER_ARTIST = None # e.g., 60 to cap for testing; None = all
HEADERS = {"User-Agent": "Mozilla/5.0 (compatible; ADS509-coursework/1.0)"}
LYRICS_DIR.mkdir(parents=True, exist_ok=True)

# sanity check
print("Twitter dir exists:", TWITTER_DIR.exists())
print("Who file exists :", TWITTER_FILES['the_who'].exists(), ">", TWITTER_FILES['the_who'])
print("Rush file exists :", TWITTER_FILES['rush'].exists(), ">", TWITTER_FILES['rush'])

Twitter dir exists: True
Who file exists : True -> C:\Users\gabed\ADS-509-Assignment-2.1\data\M1 Results\twitter\The_Who_followers_data.txt
Rush file exists : True -> C:\Users\gabed\ADS-509-Assignment-2.1\data\M1 Results\twitter\Rush_followers_data.txt
```

Helper Code Lyrics scraper + robust Twitter loader

```
In [3]: from typing import Optional, List

def polite_get(
    url: str,
    tries: int = 3,
    low: float = 1.6,
    high: float = 3.2
) -> Optional[str]:
    """
    Request a URL with retries and random back-off to avoid hammering the server.
```

```

Returns the HTML text or None if all tries fail.
"""
for _ in range(tries):
    try:
        r = requests.get(url, headers=HEADERS, timeout=15)
        if r.status_code == 200:
            return r.text
        except requests.RequestException:
            pass
    time.sleep(low + random.random() * (high - low))
return None

def extract_song_links_any(artist_index_url: str) -> List[str]:
    """
    Grab every <a> tag on the artist page whose href contains '/lyrics/'.
    Returns a list of absolute URLs to individual song pages.
    """
    html = polite_get(artist_index_url)
    if not html:
        return []
    soup = BeautifulSoup(html, "html.parser")
    return [
        urljoin(artist_index_url, a.get("href"))
        for a in soup.find_all("a", href=True)
        if "/lyrics/" in a["href"]
    ]

def parse_lyrics_page(song_url: str) -> (Optional[str], Optional[str]):
    """
    Download a song page and return (title, cleaned_lyrics).
    """
    html = polite_get(song_url)
    if not html:
        return None, None
    soup = BeautifulSoup(html, "html.parser")

    # Extract a reasonable title
    title = None
    if soup.title and soup.title.text:
        t = re.sub(r"\s*\\s*AZLyrics\.com\s*$", "", soup.title.text.strip())
        title = re.sub(r"\s*-\s*.*lyrics?$", "", t, flags=re.I).strip()

    # Find the largest block of text that looks like Lyrics
    longest = ""
    for div in soup.find_all("div"):
        if div.find(["script", "style"]) is None:
            txt = div.get_text("\n", strip=True)
            if txt and len(txt.splitlines()) > 5 and len(txt) > len(longest):
                longest = txt

    if not longest:
        return title or "Unknown Title", None

    lyr = re.sub(r"\r\n?", "\n", longest)
    lyr = re.sub(r"\n{3,}", "\n\n", lyr).strip()
    junk = [
        r"^if you found mistakes.*$", r"^submit corrections.*$", r"^writer\s\):.*$",
        r"^azlyrics.*$", r"^.*?thanks to .* for.*$", r"^album:.*$"
    ]
    cleaned = [
        ln for ln in lyr.splitlines()
        if not any(re.search(p, ln, flags=re.I) for p in junk)
    ]
    return (title or "Unknown Title"), ("\n".join(cleaned).strip() or None)

def save_lyrics(artist_key: str, title: str, body: str) -> Path:
    """
    Save a single song's lyrics as a text file under LYRICS_DIR/artist_key.
    """
    folder = LYRICS_DIR / artist_key
    folder.mkdir(parents=True, exist_ok=True)
    fname = slugify(title, lowercase=True, max_length=120) or "untitled"

```

```

path = folder / f"{fname}.txt"
with open(path, "w", encoding="utf-8") as f:
    f.write(title.strip())
    f.write("\n\n")
    f.write(body.strip())
return path

def read_lyrics_corpus(artist: str) -> List[str]:
    """
    Read all saved lyric files for a given artist into a list of strings.
    """
    texts = []
    for fp in sorted((LYRICS_DIR / artist).glob("*.txt")):
        with open(fp, "r", encoding="utf-8") as f:
            lines = f.read().splitlines()
            body = "\n".join(lines[2:]) if len(lines) >= 3 else "\n".join(lines)
            if body.strip():
                texts.append(body)
    return texts

def load_descriptions_from_file(fp: Path, sample_rows: int = 2000) -> List[str]:
    """
    Load Twitter follower descriptions from a large CSV/TSV/JSONL/text file.
    Tries CSV first, then JSON lines, then plain text.
    """
    try:
        sample = pd.read_csv(fp, sep=None, engine="python", nrows=sample_rows,
                             encoding="utf-8", on_bad_lines="skip")
        wanted = [c for c in sample.columns
                  if c.lower() in {"description", "bio", "profile_description",
                                   "user_description", "text"}]

        if wanted:
            out = []
            for chunk in pd.read_csv(fp, sep=None, engine="python", chunksize=100_000,
                                     encoding="utf-8", on_bad_lines="skip", usecols=wanted):
                for c in wanted:
                    if c in chunk.columns:
                        out.extend([str(x) for x in chunk[c].dropna().tolist()])
                        break

            if out:
                return out
    except Exception:
        pass

    # JSON lines fallback
    try:
        out = []
        with open(fp, "r", encoding="utf-8", errors="ignore") as f:
            for line in f:
                line = line.strip()
                if not line or line.startswith(("#", "///")):
                    continue
                try:
                    obj = json.loads(line)
                    for key in ["description", "bio", "profile_description",
                               "user_description", "text"]:
                        if isinstance(obj, dict) and key in obj and obj[key]:
                            out.append(str(obj[key]))
                            break
                except json.JSONDecodeError:
                    continue

        if out:
            return out
    except Exception:
        pass

    # Plain text fallback
    with open(fp, "r", encoding="utf-8", errors="ignore") as f:
        return [ln.strip() for ln in f if ln.strip()]

```

Read lyrics into memory

```
In [4]: lyrics_text = {
        "the_who": read_lyrics_corpus("the_who"),
        "rush":    read_lyrics_corpus("Rush")
    }

    twitter = {
        "the_who": load_descriptions_from_file(TWITTER_FILES["the_who"]),
        "rush":    load_descriptions_from_file(TWITTER_FILES["rush"])
    }

    print(f"The Who lyrics loaded: {len(lyrics_text['the_who'])}")
    print(f"Rush lyrics loaded:    {len(lyrics_text['rush'])}")
```

The Who lyrics loaded: 34
 Rush lyrics loaded: 316

Normalize and tokenize lyrics + twitter descriptions

```
In [6]: import re
        from nltk.tokenize import word_tokenize

        # SMALL sample for quick runs (bump later for final results)
        TW_SAMPLE = 25_000 # use None for ALL rows later

        # sanity check your inputs exist
        assert "the_who" in lyrics_text and "rush" in lyrics_text, "Run the 'Read lyrics into memory' cell."
        assert "the_who" in twitter and "rush" in twitter,      "Reload twitter descriptions into `twitter`."

        # choose a small slice of twitter to speed things up
        who_tw = twitter["the_who"][:TW_SAMPLE] if TW_SAMPLE else twitter["the_who"]
        rush_tw = twitter["rush"][:TW_SAMPLE]   if TW_SAMPLE else twitter["rush"]

        print("Input sizes:",
              "who lyrics:", len(lyrics_text["the_who"]),
              "rush lyrics:", len(lyrics_text["rush"]),
              "who twitter:", len(who_tw),
              "rush twitter:", len(rush_tw))

        EMOJI_RANGE = r"\U0001F300-\U0001FAFF"

        def normalize_lyrics(s: str) -> str:
            s = s.lower()
            s = re.sub(r"http\S+", "", s)
            s = re.sub(r"^[a-z0-9\s']+", " ", s)
            return re.sub(r"\s+", " ", s).strip()

        def normalize_twitter(s: str) -> str:
            s = s.lower()
            s = re.sub(r"http\S+", "", s)
            s = re.sub(fr"^[^\s#\/\._{EMOJI_RANGE}]", " ", s) # keep hashtags/_.- and emojis
            return re.sub(r"\s+", " ", s).strip()

        def tokenize_with_progress(texts, normalizer, label="", every=10_000):
            out = []
            n = len(texts)
            for i, t in enumerate(texts, 1):
                out.extend(tok for tok in word_tokenize(normalizer(t)) if tok.strip())
                if label and (i % every == 0 or i == n):
                    print(f"{label}: {i:,}/{n:,}")
            return out

        print("Tokenizing...")
        tokens = {
            "the_who_lyrics": tokenize_with_progress(lyrics_text["the_who"], normalize_lyrics, "who lyrics", 200),
            "rush_lyrics":    tokenize_with_progress(lyrics_text["rush"],    normalize_lyrics, "rush lyrics", 500),
            "the_who_twitter": tokenize_with_progress(who_tw,                normalize_twitter, "who twitter"),
            "rush_twitter":    tokenize_with_progress(rush_tw,                normalize_twitter, "rush twitter"),
        }

        print("Token counts:", {k: len(v) for k, v in tokens.items()})
```

Input sizes: who lyrics: 34 rush lyrics: 316 who twitter: 25000 rush twitter: 25000
Tokenizing...
who lyrics: 34/34
rush lyrics: 316/316
who twitter: 10,000/25,000
who twitter: 20,000/25,000
who twitter: 25,000/25,000
rush twitter: 10,000/25,000
rush twitter: 20,000/25,000
rush twitter: 25,000/25,000
Token counts: {'the_who_lyrics': 10589, 'rush_lyrics': 73844, 'the_who_twitter': 383919, 'rush_twitter': 336155}

Descriptive statistics for the lyrics corpora

```
In [7]: def describe(token_list):
        N = len(token_list)          # total tokens
        V = len(set(token_list))     # vocabulary size
        return {
            "tokens": N,
            "vocabulary": V,
            "type_token_ratio": V / N if N else 0,
            "top20": Counter(token_list).most_common(20)
        }

        # calculate for lyrics only
        lyrics_stats = {
            artist: describe(tokens[f"{artist}_lyrics"])
            for artist in ["the_who", "rush"]
        }

        # show summary table
        import pandas as pd
        summary = pd.DataFrame([
            {
                "artist": a,
                "tokens": v["tokens"],
                "vocabulary": v["vocabulary"],
                "type_token_ratio": round(v["type_token_ratio"], 3),
                "top20": v["top20"]
            }
            for a, v in lyrics_stats.items()
        ])
        summary
```

```
Out[7]:
```

	artist	tokens	vocabulary	type_token_ratio	top20
0	the_who	10589	1221	0.115	[(the, 430), (you, 428), (i, 394), (to, 296), ...]
1	rush	73844	3681	0.050	[(i, 3566), (you, 3261), (the, 2395), (to, 166...]

Unique words in each corpus

```
In [8]: # create sets of tokens for easy set operations
        token_sets = {name: set(tok_list) for name, tok_list in tokens.items()}

        unique_words = {}
        for name, current_set in token_sets.items():
            others = set().union(*[s for k, s in token_sets.items() if k != name])
            unique_words[name] = sorted(current_set - others)

        # quick check: number of unique words per corpus
        for name, uw in unique_words.items():
            print(f"{name}: {len(uw)} unique words")

        # preview a small sample from each
        for name, uw in unique_words.items():
            print(f"\n{name} - first 25 unique words:")
            print(uw[:25])
```


[illegible][illegible]

